

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import random
```

```
In [2]: df = pd.read_csv("bank-full.csv")
df
```

```
Out[2]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may
...
45206	51	technician	married	tertiary	no	825	no	no	cellular	17	may
45207	71	retired	divorced	primary	no	1729	no	no	cellular	17	may
45208	72	retired	married	secondary	no	5715	no	no	cellular	17	may
45209	57	blue-collar	married	secondary	no	668	no	no	telephone	17	may
45210	37	entrepreneur	married	secondary	no	2971	no	no	cellular	17	may

45211 rows × 12 columns



```
In [3]: df.head()
```

```
Out[3]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may



```
In [4]: df.shape
```

```
Out[4]: (45211, 12)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
```

```
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null    int64
1   job          45211 non-null    object
2   marital      45211 non-null    object
3   education    45211 non-null    object
4   default      45211 non-null    object
5   balance      45211 non-null    int64
6   housing      45211 non-null    object
7   loan         45211 non-null    object
8   contact      45211 non-null    object
9   day          45211 non-null    int64
10  month        45211 non-null    object
11  duration     45211 non-null    int64
12  campaign     45211 non-null    int64
13  pdays        45211 non-null    int64
14  previous     45211 non-null    int64
15  poutcome     45211 non-null    object
16  Target       45211 non-null    object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

In [6]: `df.isna().sum()`

```
Out[6]: age      0
job      0
marital   0
education 0
default   0
balance   0
housing   0
loan      0
contact   0
day       0
month     0
duration  0
campaign  0
pdays    0
previous  0
poutcome  0
Target    0
dtype: int64
```

In [7]: `df.describe(include = 'all')`

Out[7]:

	age	job	marital	education	default	balance	housing	loan	contact
count	45211.000000	45211	45211	45211	45211	45211.000000	45211	45211	45211
unique	NaN	12	3	4	2	NaN	2	2	3
top	NaN	blue-collar	married	secondary	no	NaN	yes	no	cellular
freq	NaN	9732	27214	23202	44396	NaN	25130	37967	29285
mean	40.936210	NaN	NaN	NaN	NaN	1362.272058	NaN	NaN	NaN
std	10.618762	NaN	NaN	NaN	NaN	3044.765829	NaN	NaN	NaN
min	18.000000	NaN	NaN	NaN	NaN	-8019.000000	NaN	NaN	NaN
25%	33.000000	NaN	NaN	NaN	NaN	72.000000	NaN	NaN	NaN
50%	39.000000	NaN	NaN	NaN	NaN	448.000000	NaN	NaN	NaN

	age	job	marital	education	default	balance	housing	loan	contact
75%	48.000000	NaN	NaN	NaN	NaN	1428.000000	NaN	NaN	NaN
max	95.000000	NaN	NaN	NaN	NaN	102127.000000	NaN	NaN	NaN

In [8]: `df.dtypes`

Out[8]:

```

age          int64
job          object
marital      object
education    object
default      object
balance      int64
housing      object
loan         object
contact      object
day          int64
month        object
duration     int64
campaign     int64
pdays      int64
previous     int64
poutcome    object
Target      object
dtype: object

```

Manipulating the data

In [9]:

```

#Creating User Columns
df_user = pd.DataFrame(np.arange(0,len(df)), columns=['user'])
df = pd.concat([df_user, df], axis=1)

```

In [10]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 18 columns):
#   Column      Non-Null Count  Dtype
---  -
0   user        45211 non-null  int32
1   age         45211 non-null  int64
2   job         45211 non-null  object
3   marital     45211 non-null  object
4   education   45211 non-null  object
5   default     45211 non-null  object
6   balance     45211 non-null  int64
7   housing     45211 non-null  object
8   loan        45211 non-null  object
9   contact     45211 non-null  object
10  day         45211 non-null  int64
11  month       45211 non-null  object
12  duration    45211 non-null  int64
13  campaign    45211 non-null  int64
14  pdays       45211 non-null  int64
15  previous    45211 non-null  int64
16  poutcome    45211 non-null  object
17  Target      45211 non-null  object
dtypes: int32(1), int64(7), object(10)
memory usage: 6.0+ MB

```

```
In [11]: df.columns.values
```

```
Out[11]: array(['user', 'age', 'job', 'marital', 'education', 'default', 'balance',
        'housing', 'loan', 'contact', 'day', 'month', 'duration',
        'campaign', 'pdays', 'previous', 'poutcome', 'Target'],
        dtype=object)
```

```
In [12]: df.groupby('Target').mean()
```

```
Out[12]:
```

	user	age	balance	day	duration	campaign	pdays	previous
Target								
no	21197.503081	40.838986	1303.714969	15.892290	221.182806	2.846350	36.421372	0.502154
yes	33228.953867	41.670070	1804.267915	15.158253	537.294574	2.141047	68.702968	1.170354

```
In [13]: df['Target'].value_counts()
```


```
Out[13]: no      39922
        yes      5289
        Name: Target, dtype: int64
```

```
In [14]: countNo = len(df[df.Target == 'no'])
        countYes = len(df[df.Target == 'yes'])
        print('Percentage of "No": {:.3f}%'.format((countNo/(len(df.Target))*100)))
        print('Percentage of "Yes": {:.3f}%'.format((countYes/(len(df.Target))*100)))
```

```
Percentage of "No": 88.302%
Percentage of "Yes": 11.698%
```

```
In [15]: #Verifying null values
        sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap='viridis')
```

```
Out[15]: <AxesSubplot:>
```



```
In [16]: df.isna().any()
```

```
Out[16]: user      False
        age      False
```

```

job            False
marital        False
education      False
default        False
balance        False
housing        False
loan           False
contact        False
day            False
month          False
duration       False
campaign       False
pdays         False
previous       False
poutcome       False
Target         False
dtype: bool

```

```
In [17]: df.isna().sum()
```

```

Out[17]: user            0
age              0
job              0
marital          0
education        0
default          0
balance          0
housing          0
loan             0
contact          0
day              0
month            0
duration         0
campaign         0
pdays          0
previous         0
poutcome         0
Target           0
dtype: int64

```

```

In [18]: #Define X and y
X = df.drop(['Target','user','job','marital','education','contact',
            'housing','loan','day','month','poutcome'], axis=1)
y = df['Target']

```

```

In [19]: X = pd.get_dummies(X)
y = pd.get_dummies(y)

```

```

In [20]: X.columns
X = X.drop(['default_no'], axis=1)
X = X.rename(columns = {'default_yes': 'default'})
y.columns
y = y.drop(['yes'], axis=1)
y = y.rename(columns= {'no': 'y'})

```

visualizing data

```

In [21]: #Age group
bins = range(0, 100, 10)
ax = sns.distplot(df.age[df.Target=='yes'],

```


```

        color='red', kde=False, bins=bins, label='Have Subscribed')
sns.distplot(df.age[df.Target=='no'],
             ax=ax, # Overplots on first plot
             color='blue', kde=False, bins=bins, label="Haven't Subscribed")
plt.legend()
plt.show()

```

C:\Users\Admin\sachin\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)




In [22]:

```

#Age
pd.crosstab(df.age,df.Target).plot(kind="bar",figsize=(20,6))
plt.title('Client Subscribed Frequency for Ages')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()

```

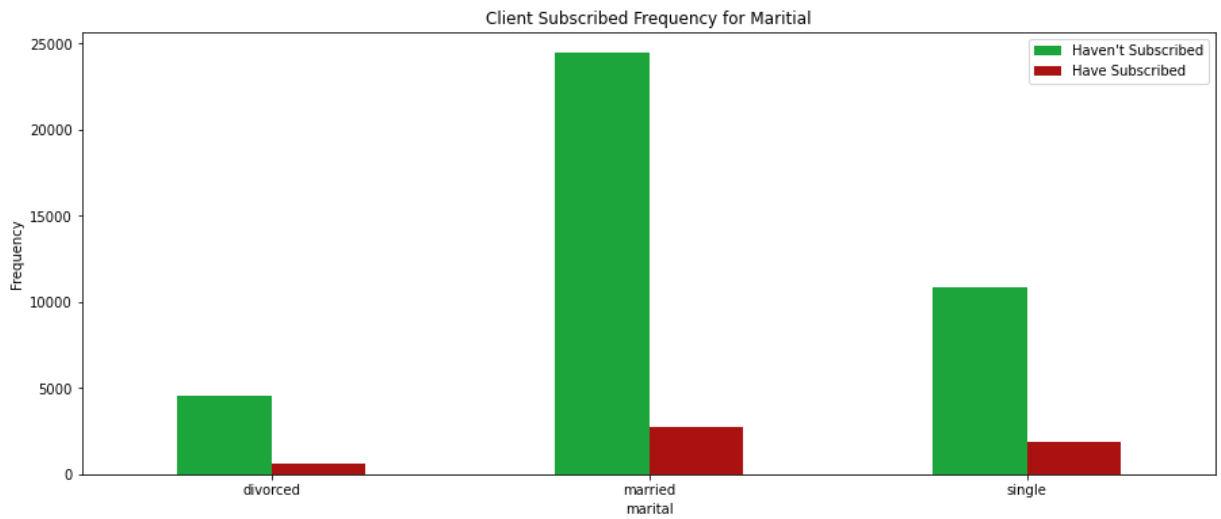


In [23]:


```

pd.crosstab(df.marital,df.Target).plot(kind="bar",figsize=(15,6),color=['#1CA53B','#
plt.title('Client Subscribed Frequency for Marital')
plt.xlabel('marital')
plt.xticks(rotation=0)
plt.legend(["Haven't Subscribed", "Have Subscribed"])
plt.ylabel('Frequency')
plt.show()

```

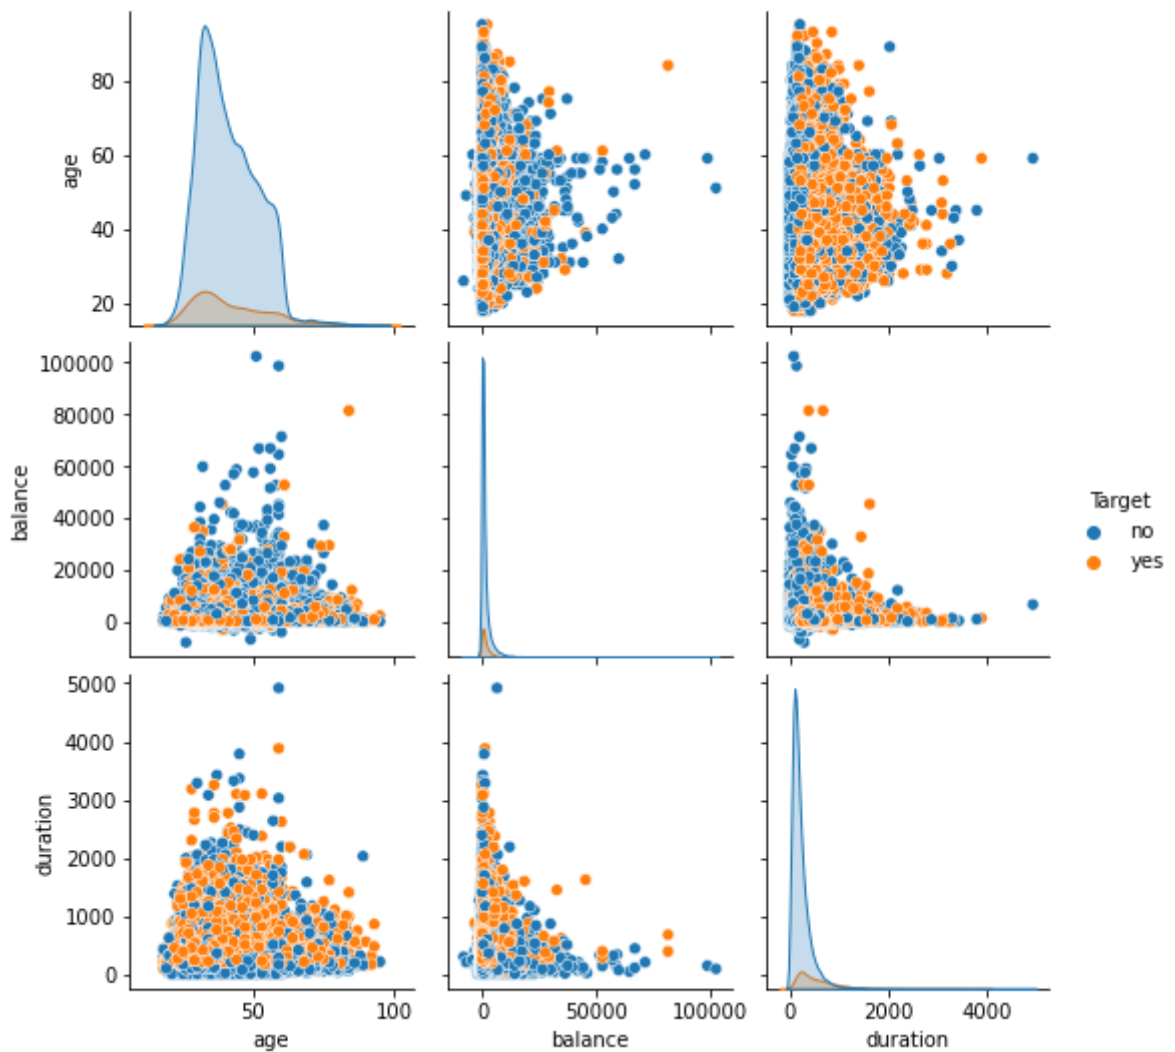


```
In [24]: plt.scatter(x=df.age[df.Target=='yes'], y=df.duration[(df.Target=='yes')], c="red")
plt.scatter(x=df.age[df.Target=='no'], y=df.duration[(df.Target=='no')])
plt.legend(["Have Subscribed", "Haven't Subscribed"])
plt.xlabel("Age")
plt.ylabel("Duration")
plt.show()
```




```
In [25]: sns.pairplot(data=df, hue='Target', vars= ['age', 'balance', 'duration'])
```

```
Out[25]: <seaborn.axisgrid.PairGrid at 0x1c9abf612e0>
```




```
In [26]: sns.countplot(x='Target', data=df, label='Count')
```

```
Out[26]: <AxesSubplot:xlabel='Target', ylabel='count'>
```



```
In [27]: sns.scatterplot(x='age', y='balance', hue='Target', data=df)
```

```
Out[27]: <AxesSubplot:xlabel='age', ylabel='balance'>
```

```
In [28]: plt.figure(figsize=(20,10))
sns.heatmap(data=df.corr(), annot=True, cmap='viridis')
```


Out[28]: <AxesSubplot:>



```
In [29]: import warnings
warnings.filterwarnings("ignore")
```

```
In [30]: sns.distplot(df.balance, bins = 20)
```

Out[30]: <AxesSubplot:xlabel='balance', ylabel='Density'>

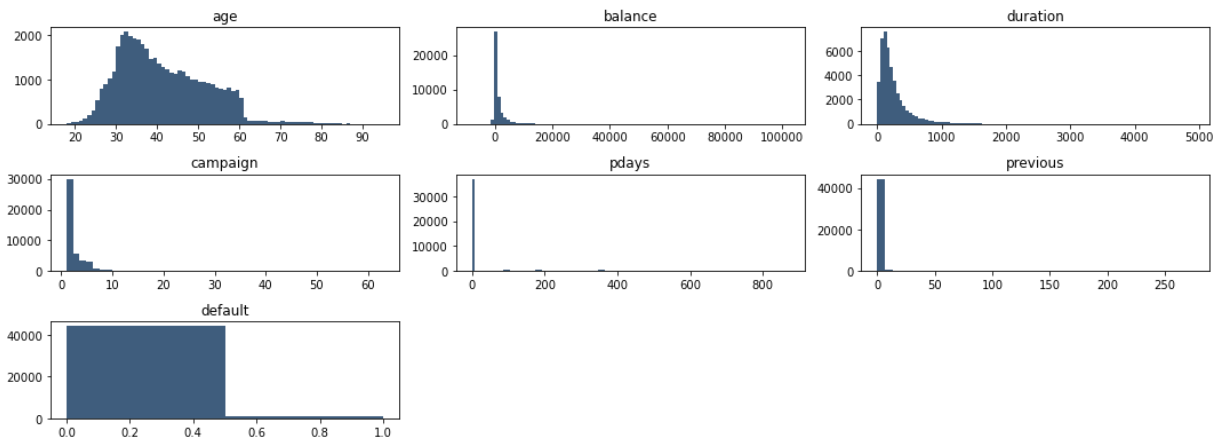


```
In [31]: df2 = X
fig = plt.figure(figsize=(15, 12))
plt.suptitle('Histograms of Numerical Columns', fontsize=20)
for i in range(df2.shape[1]):
    plt.subplot(6, 3, i + 1)
    f = plt.gca()
    f.set_title(df2.columns.values[i])

    vals = np.size(df2.iloc[:, i].unique())
    if vals >= 100:
        vals = 100


    plt.hist(df2.iloc[:, i], bins=vals, color='#3F5D7D')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

Histograms of Numerical Columns



```
In [32]: ## Correlation with independent Variable
df2.corrwith(y.y).plot.bar(
    figsize = (10, 10), title = "Correlation with Y", fontsize = 15,
    rot = 45, grid = True)
```

```
Out[32]: <AxesSubplot:title={'center':'Correlation with Y'}>
```



```
In [33]: # Compute the correlation matrix
corr = df2.corr()


# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(10, 10))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

Out[33]: <AxesSubplot:>




In [34]:

```
## Pie Plots
df.columns
df2 = df[['Target','job','marital', 'education', 'default', 'housing','loan', 'contact',
          'month', 'outcome'
          ]]
fig = plt.figure(figsize=(15, 12))
plt.suptitle('Pie Chart Distributions', fontsize=20)
for i in range(1, df2.shape[1] + 1):
    plt.subplot(6, 3, i)
    f = plt.gca()
    f.axes.get_yaxis().set_visible(False)
    f.set_title(df2.columns.values[i - 1])

    values = df2.iloc[:, i - 1].value_counts(normalize = True).values
    index = df2.iloc[:, i - 1].value_counts(normalize = True).index
    plt.pie(values, labels = index, autopct='%1.1f%%')
    plt.axis('equal')
fig.tight_layout(rect=[0, 0.03, 1, 0.95])
```

Pie Chart Distributions



Splitting data into training and testing data

```
In [35]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y,
```

```
In [36]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(36168, 7) (9043, 7) (36168, 1) (9043, 1)
```

```
In [37]: ## Balance the data
y_train['y'].value_counts()
```

```
Out[37]: 1    31937
         0     4231
         Name: y, dtype: int64
```

```
In [38]: pos_index = y_train[y_train.values == 1].index
neg_index = y_train[y_train.values == 0].index

if len(pos_index) > len(neg_index):
    higher = pos_index
    lower = neg_index
else:
    higher = neg_index
    lower = pos_index

random.seed(0)
higher = np.random.choice(higher, size=len(lower))
lower = np.asarray(lower)
new_indexes = np.concatenate((lower, higher))

X_train = X_train.loc[new_indexes]
y_train = y_train.loc[new_indexes]
```

```
In [39]: y_train['y'].value_counts()
```

```
0    4231
```

```
Out[39]: 1    4231
         Name: y, dtype: int64
```

Feature Scaling

```
In [40]: from sklearn.preprocessing import StandardScaler
         sc = StandardScaler()
         X_train2 = pd.DataFrame(sc.fit_transform(X_train))
         X_test2 = pd.DataFrame(sc.transform(X_test))
         X_train2.columns = X_train.columns.values
         X_test2.columns = X_test.columns.values
         X_train2.index = X_train.index.values
         X_test2.index = X_test.index.values
         X_train = X_train2
         X_test = X_test2
```

Model Buliding

Comparing models

```
In [42]: ## Logistic Regression
         from sklearn.linear_model import LogisticRegression
         classifier = LogisticRegression(random_state = 0, penalty = 'l2')
         classifier.fit(X_train, y_train)

         # Predicting Test Set
         y_pred = classifier.predict(X_test)
         from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score
         acc = accuracy_score(y_test, y_pred)
         prec = precision_score(y_test, y_pred)
         rec = recall_score(y_test, y_pred)
         f1 = f1_score(y_test, y_pred)

         results = pd.DataFrame([['Logistic Regression (Lasso)', acc, prec, rec, f1]],
                                columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
```

```
In [48]: results
```

```
Out[48]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335

```
In [50]: ## SVM (Linear)
         from sklearn.svm import SVC
         classifier = SVC(random_state = 0, kernel = 'linear', probability= True)
         classifier.fit(X_train, y_train)
```

```
# Predicting Test Set
y_pred = classifier.predict(X_test)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

model_results = pd.DataFrame([['SVM (Linear)', acc, prec, rec, f1]],
                             columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

results = results.append(model_results, ignore_index = True)
```

In [51]: results

Out[51]:

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335
6	SVM (Linear)	0.797965	0.952661	0.811522	0.876446

In [44]:

```
## SVM (rbf)
from sklearn.svm import SVC
classifier = SVC(random_state = 0, kernel = 'rbf', probability= True)
classifier.fit(X_train, y_train)

# Predicting Test Set
y_pred = classifier.predict(X_test)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

model_results = pd.DataFrame([['SVM (RBF)', acc, prec, rec, f1]],
                             columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

results = results.append(model_results, ignore_index = True)
```

In [52]: results

Out[52]:

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335

	Model	Accuracy	Precision	Recall	F1 Score
6	SVM (Linear)	0.797965	0.952661	0.811522	0.876446

```
In [45]: ## Naive Bayes
from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train, y_train)

# Predicting Test Set
y_pred = classifier.predict(X_test)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

model_results = pd.DataFrame([['Naive Bayes (Gaussian)', acc, prec, rec, f1]],
                              columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

results = results.append(model_results, ignore_index = True)
```

```
In [53]: results
```

```
Out[53]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335
6	SVM (Linear)	0.797965	0.952661	0.811522	0.876446

```
In [46]: # Decision Tree
from sklearn.tree import DecisionTreeClassifier
classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)
classifier.fit(X_train, y_train)

#Predicting the best set result
y_pred = classifier.predict(X_test)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

model_results = pd.DataFrame([['Decision Tree', acc, prec, rec, f1]],
                              columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

results = results.append(model_results, ignore_index = True)
```

```
In [54]: results
```

```
Out[54]:
```

	Model	Accuracy	Precision	Recall	F1 Score
--	-------	----------	-----------	--------	----------

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335
6	SVM (Linear)	0.797965	0.952661	0.811522	0.876446

```
In [47]: ## Random Forest Gini (n=100)
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(random_state = 0, n_estimators = 100,
                                  criterion = 'gini')
classifier.fit(X_train, y_train)

# Predicting Test Set
y_pred = classifier.predict(X_test)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

model_results = pd.DataFrame([['Random Forest Gini (n=100)', acc, prec, rec, f1]],
                             columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

results = results.append(model_results, ignore_index = True)
```

```
In [55]: results
```

```
Out[55]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335
6	SVM (Linear)	0.797965	0.952661	0.811522	0.876446

Applying k-fold validation

```
In [56]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator=classifier, X=X_train, y=y_train, cv=10)
accuracies.mean()
accuracies.std()
```


```
Out[56]: 0.009004867531703015
```

In [57]: `print("SVM (Linear) Accuracy: %0.3f (+/- %0.3f)" % (accuracies.mean(), accuracies.st`

SVM (Linear) Accuracy: 0.751 (+/- 0.018)

In [58]: `##### Confusion Matrix
cm = confusion_matrix(y_test, y_pred) # rows = truth, cols = prediction
df_cm = pd.DataFrame(cm, index = (0, 1), columns = (0, 1))
plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, fmt='g')
print("Test Data Accuracy: %0.4f" % accuracy_score(y_test, y_pred))`

Test Data Accuracy: 0.7980



In [59]: `#Plotting Cumulative Accuracy Profile (CAP)
y_pred_proba = classifier.predict_proba(X=X_test)
import matplotlib.pyplot as plt
from scipy import integrate
def capcurve(y_values, y_preds_proba):
 num_pos_obs = np.sum(y_values)
 num_count = len(y_values)
 rate_pos_obs = float(num_pos_obs) / float(num_count)
 ideal = pd.DataFrame({'x':[0,rate_pos_obs,1], 'y':[0,1,1]})
 xx = np.arange(num_count) / float(num_count - 1)

 y_cap = np.c_[y_values,y_preds_proba]
 y_cap_df_s = pd.DataFrame(data=y_cap)
 y_cap_df_s = y_cap_df_s.sort_values([1], ascending=False).reset_index(level = y_

 print(y_cap_df_s.head(20))

 yy = np.cumsum(y_cap_df_s[0]) / float(num_pos_obs)
 yy = np.append([0], yy[0:num_count-1]) #add the first curve point (0,0) : for xx

 percent = 0.5
 row_index = int(np.trunc(num_count * percent))`

```

val_y1 = yy[row_index]
val_y2 = yy[row_index+1]
if val_y1 == val_y2:
    val = val_y1*1.0
else:
    val_x1 = xx[row_index]
    val_x2 = xx[row_index+1]
    val = val_y1 + ((val_x2 - percent)/(val_x2 - val_x1))*(val_y2 - val_y1)

sigma_ideal = 1 * xx[num_pos_obs - 1] / 2 + (xx[num_count - 1] - xx[num_pos_obs
sigma_model = integrate.simps(yy,xx)
sigma_random = integrate.simps(xx,xx)

ar_value = (sigma_model - sigma_random) / (sigma_ideal - sigma_random)

fig, ax = plt.subplots(nrows = 1, ncols = 1)
ax.plot(ideal['x'],ideal['y'], color='grey', label='Perfect Model')
ax.plot(xx,yy, color='red', label='User Model')
ax.plot(xx,xx, color='blue', label='Random Model')
ax.plot([percent, percent], [0.0, val], color='green', linestyle='--', linewidth=1, la
ax.plot([0, percent], [val, val], color='green', linestyle='--', linewidth=1, la


plt.xlim(0, 1.02)
plt.ylim(0, 1.25)
plt.title("CAP Curve - a_r value =" +str(ar_value))
plt.xlabel('% of the data')
plt.ylabel('% of positive obs')
plt.legend()

```

In [60]:

```
capcurve(y_test,y_pred_proba[:,1])
```

	0	1
0	1.0	0.999993
1	1.0	0.995807
2	1.0	0.991640
3	1.0	0.990919
4	1.0	0.990914
5	1.0	0.989523
6	1.0	0.987459
7	1.0	0.986814
8	1.0	0.986797
9	1.0	0.986063
10	1.0	0.986025
11	1.0	0.985244
12	1.0	0.985150
13	1.0	0.983236
14	1.0	0.982235
15	1.0	0.981915
16	1.0	0.980047
17	1.0	0.978314
18	1.0	0.978290
19	1.0	0.977016



```
In [61]: # Analyzing Coefficients
pd.concat([pd.DataFrame(X_train.columns, columns = ["features"]),
           pd.DataFrame(np.transpose(classifier.coef_), columns = ["coef"])
           ],axis = 1)
```

```
Out[61]:
```

	features	coef
0	age	-0.066575
1	balance	-0.193843
2	duration	-1.528572
3	campaign	0.204395
4	pdays	-0.157947
5	previous	-0.425870
6	default	0.026428

```
In [62]: ##### Recursive Feature Elimination
from sklearn.feature_selection import RFE
from sklearn.svm import SVC

# Model to Test
classifier = SVC(random_state = 0, kernel = 'linear', probability= True)

# Select Best X Features
rfe = RFE(classifier, n_features_to_select=None)
rfe = rfe.fit(X_train, y_train)
```

```
In [63]: # summarize the selection of the attributes
print(rfe.support_)
print(rfe.ranking_)
```

```
[False False  True  True False  True False]
[4 2 1 1 3 1 5]
```

```
In [64]: X_train.columns[rfe.support_]
```

```
Out[64]: Index(['duration', 'campaign', 'previous'], dtype='object')
```

```

In [65]: # New Correlation Matrix
sns.set(style="white")

# Compute the correlation matrix
corr = X_train[X_train.columns[rfe.support_]].corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True


# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(18, 15))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

```

Out[65]: <AxesSubplot:>



```

In [66]: # Fitting Model to the Training Set
classifier = SVC(random_state = 0, kernel = 'linear', probability= True)
classifier.fit(X_train[X_train.columns[rfe.support_]], y_train)

# Predicting Test Set
y_pred = classifier.predict(X_test[X_train.columns[rfe.support_]])

```

```
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

model_results = pd.DataFrame([['SVM RFE (Linear)', acc, prec, rec, f1]],
                              columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

results = results.append(model_results, ignore_index = True)
```

In [67]: results


Out[67]:

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression (Lasso)	0.794095	0.952952	0.806637	0.873711
1	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
2	SVM (RBF)	0.782705	0.966956	0.780589	0.863835
3	Naive Bayes (Gaussian)	0.773195	0.948594	0.785723	0.859511
4	Decision Tree	0.719894	0.954940	0.716594	0.818774
5	Random Forest Gini (n=100)	0.768771	0.969866	0.761803	0.853335
6	SVM (Linear)	0.797965	0.952661	0.811522	0.876446
7	SVM RFE (Linear)	0.803937	0.951060	0.820163	0.880775

In [68]:

```
# Evaluating Results
#Making the confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test,y_pred)
sns.heatmap(data=cm, annot=True)
```

Out[68]: <AxesSubplot:>



In [69]:

```
#Making the classification report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.33	0.68	0.45	1058
1	0.95	0.82	0.88	7985

accuracy			0.80	9043
macro avg	0.64	0.75	0.66	9043
weighted avg	0.88	0.80	0.83	9043

```
In [70]: # Applying k-Fold Cross Validation (RFE)
from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = classifier,
                              X = X_train[X_train.columns[rfe.support_]],
                              y = y_train, cv = 10)
```

```
In [72]: print("SVM RFE (Linear) Accuracy: %0.3f (+/- %0.3f)" % (accuracies.mean(), accuracies.std())

SVM RFE (Linear) Accuracy: 0.747 (+/- 0.020)
```


```
In [73]: # Analyzing Coefficients
pd.concat([pd.DataFrame(X_train[X_train.columns[rfe.support_]].columns, columns = ["feature"]),
           pd.DataFrame(np.transpose(classifier.coef_), columns = ["coef"])
          ],axis = 1)
```

```
Out[73]:
```

	features	coef
0	duration	-1.539859
1	campaign	0.214335
2	previous	-0.574785

```
In [74]: #CAP Curve
y_pred_proba = classifier.predict_proba(X=X_test[X_train.columns[rfe.support_]])
capcurve(y_test,y_pred_proba[:,1])
```

	0	1
0	1.0	0.999990
1	1.0	0.995396
2	1.0	0.990650
3	1.0	0.990570
4	1.0	0.990336
5	1.0	0.987674
6	1.0	0.986908
7	1.0	0.986697
8	1.0	0.986470
9	1.0	0.984798
10	1.0	0.984451
11	1.0	0.983978
12	1.0	0.983303
13	1.0	0.983208
14	1.0	0.981000
15	1.0	0.979525
16	1.0	0.977628
17	1.0	0.976670
18	1.0	0.974681
19	1.0	0.972450



```
In [75]: ### End of the Model

# Formatting Final Results
user_idenfifier = df['user']
final_results = pd.concat([y_test, user_idenfifier], axis = 1).dropna()
final_results['predicted'] = y_pred
final_results = final_results[['user', 'y', 'predicted']].reset_index(drop=True)
```

```
In [76]: final_results.head()
```

Out[76]:

	user	y	predicted
0	5	1.0	1
1	11	1.0	1
2	20	1.0	1
3	21	1.0	1
4	26	1.0	1

```
In [ ]:
```